





Predicting House Prices on Kaggle

Eng Teong Cheah

Contents


Predicting House Prices on Kaggle


Kaggle

Competitions Datasets Kernels Discussion Learn ... [Sign In](#)

Kaggle is the place to do data science projects

[See how it works](#) 



Sign up with just one click:


We won't share anything without your permission

[Google](#) [Facebook](#) [Yahoo](#)

Manually create an account:

[Sign Up](#)

Kaggle



House Prices: Advanced Regression Techniques

Predict sales prices and practice feature engineering, RFs, and gradient boosting

5,012 teams · Ongoing

[Overview](#) [Data](#) [Kernels](#) [Discussion](#) [Leaderboard](#) [Rules](#) [Team](#) [My Submissions](#) [Submit Predictions](#)

Overview

Description	<h4>Start here if...</h4> <p>You have some experience with R or Python and machine learning basics. This is a perfect competition for data science students who have completed an online course in machine learning and are looking to expand their skill set before trying a featured competition.</p>
Evaluation	
Frequently Asked Questions	
Tutorials	<h4>Competition Description</h4>

<https://www.kaggle.com/c/house-prices-advanced-regression-techniques>

Accessing and Reading Data Sets



```
# If pandas is not installed, please uncomment the following line:  
# !pip install pandas  
  
%matplotlib inline  
import d2l  
from mxnet import autograd, gluon, init, nd  
from mxnet.gluon import data as gdata, loss as gloss, nn  
import numpy as np  
import pandas as pd
```

Accessing and Reading Data Sets



```
train_data = pd.read_csv(' ../data/kaggle_house_pred_train.csv')  
test_data = pd.read_csv(' ../data/kaggle_house_pred_test.csv')
```

Accessing and Reading Data Sets



```
print(train_data.shape)  
print(test_data.shape)
```

Accessing and Reading Data Sets



```
train_data.iloc[0:4, [0, 1, 2, 3, -3, -2,  
-1]]
```

	Id	MSSubClass	MSZoning	LotFrontage	SaleType	SaleCondition	SalePrice
0	1	60	RL	65.0	WD	Normal	208500
1	2	20	RL	80.0	WD	Normal	181500
2	3	60	RL	68.0	WD	Normal	223500
3	4	70	RL	60.0	WD	Abnorml	140000

Accessing and Reading Data Sets



```
all_features = pd.concat((train_data.iloc[:, 1:-1], test_data.iloc[:,  
1:]))
```

Data Preprocessing



```
numeric_features = all_features.dtypes[all_features.dtypes !=  
object].index  
all_features[numeric_features] = all_features[numeric_features].apply(  
    lambda x: (x - x.mean()) / (x.std()))  
# After standardizing the data all means vanish, hence we can set missing  
# values to 0  
all_features[numeric_features] = all_features[numeric_features].fillna(0)
```

Data Preprocessing



```
# Dummy_na=True refers to a missing value being a legal eigenvalue, and  
# creates an indicative feature for it  
all_features = pd.get_dummies(all_features, dummy_na=True)  
all_features.shape
```

Data Preprocessing



```
n_train = train_data.shape[0]
train_features = nd.array(all_features[:n_train].values)
test_features = nd.array(all_features[n_train:].values)
train_labels = nd.array(train_data.SalePrice.values).reshape((-1,
1))
```

Training



```
loss = gloss.L2Loss()
```

```
def get_net():  
    net = nn.Sequential()  
    net.add(nn.Dense(1))  
    net.initialize()  
    return net
```

Training

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (\log(y_i) - \log(\hat{y}_i))^2}$$




```
def log_rmse(net, features, labels):  
    # To further stabilize the value when the logarithm is taken, set the  
    # value less than 1 as 1  
    clipped_preds = nd.clip(net(features), 1, float('inf'))  
    rmse = nd.sqrt(2 * loss(clipped_preds.log(), labels.log()).mean())  
    return rmse.asscalar()
```

Training

```
def train(net, train_features, train_labels, test_features, test_labels,
          num_epochs, learning_rate, weight_decay, batch_size):
    train_ls, test_ls = [], []
    train_iter = gdata.DataLoader(gdata.ArrayDataset(
        train_features, train_labels), batch_size, shuffle=True)
    # The Adam optimization algorithm is used here
    trainer = gluon.Trainer(net.collect_params(), 'adam', {
        'learning_rate': learning_rate, 'wd': weight_decay})
    for epoch in range(num_epochs):
        for X, y in train_iter:
            with autograd.record():
                l = loss(net(X), y)
            l.backward()
            trainer.step(batch_size)
        train_ls.append(log_rmse(net, train_features, train_labels))
        if test_labels is not None:
            test_ls.append(log_rmse(net, test_features, test_labels))
    return train_ls, test_ls
```

k-Fold Cross-Validation



```
def get_k_fold_data(k, i, X, y):
    assert k > 1
    fold_size = X.shape[0] // k
    X_train, y_train = None, None
    for j in range(k):
        idx = slice(j * fold_size, (j + 1) *
fold_size), y_part = X[idx, :], y[idx]
        if j == i:
            X_valid, y_valid = X_part, y_part
        elif X_train is None:
            X_train, y_train = X_part, y_part
        else:
            X_train = nd.concat(X_train, X_part,
dim=0)
            y_train = nd.concat(y_train, y_part,
dim=0)
    return X_train, y_train, X_valid, y_valid
```


k-Fold Cross-Validation

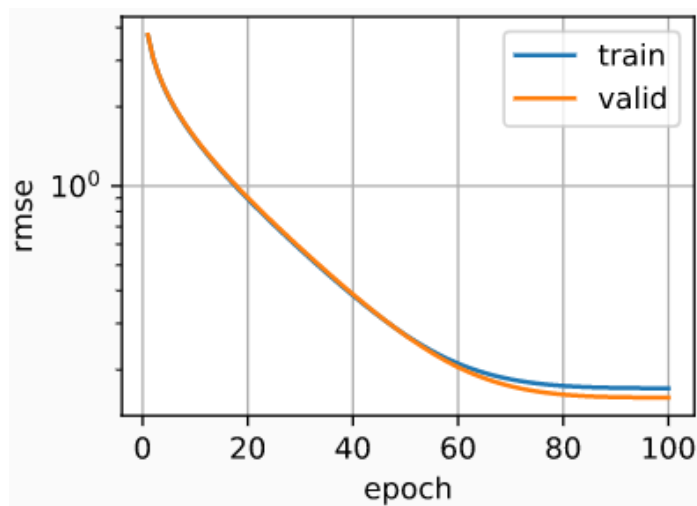
```
def k_fold(k, X_train, y_train, num_epochs,
           learning_rate, weight_decay, batch_size):
    train_l_sum, valid_l_sum = 0, 0
    for i in range(k):
        data = get_k_fold_data(k, i, X_train, y_train)
        net = get_net()
        train_ls, valid_ls = train(net, *data, num_epochs, learning_rate,
                                   weight_decay, batch_size)

        train_l_sum += train_ls[-1]
        valid_l_sum += valid_ls[-1]
        if i == 0:
            d2l.plot(list(range(1, num_epochs+1)), [train_ls, valid_ls],
                     xlabel='epoch', ylabel='rmse',
                     legend=['train', 'valid'], yscale='log')
        print('fold %d, train rmse: %f, valid rmse: %f' % (
            i, train_ls[-1], valid_ls[-1]))
    return train_l_sum / k, valid_l_sum / k
```


Model Selection



```
k, num_epochs, lr, weight_decay, batch_size = 5, 100, 5, 0, 64
train_l, valid_l = k_fold(k, train_features, train_labels, num_epochs,
                           lr, weight_decay, batch_size)
print('%d-fold validation: avg train rmse: %f, avg valid rmse: %f'
      % (k, train_l, valid_l))
```



Predict and Submit

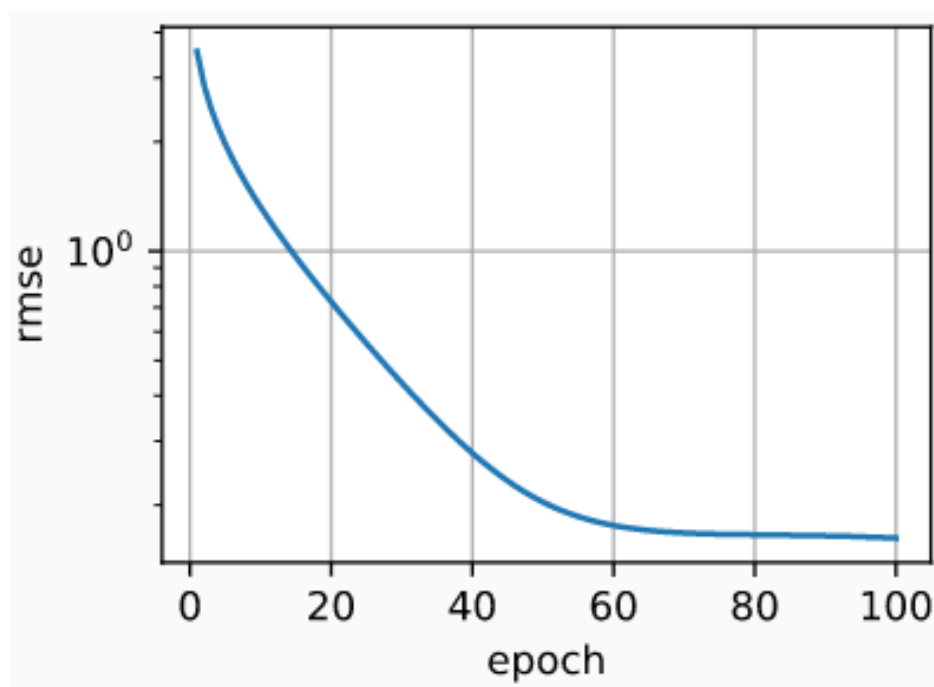


```
def train_and_pred(train_features, test_features, train_labels, test_data,
                   num_epochs, lr, weight_decay, batch_size):
    net = get_net()
    train_ls, _ = train(net, train_features, train_labels, None, None,
                        num_epochs, lr, weight_decay, batch_size)
    d2l.semilogy(range(1, num_epochs + 1), train_ls, 'epochs', 'rmse')
    print('train rmse %f' % train_ls[-1])
    preds = net(test_features).asnumpy()
    test_data['SalePrice'] = pd.Series(preds.reshape(1, -1)[0])
    submission = pd.concat([test_data['Id'], test_data['SalePrice']],
                           axis=1)
    submission.to_csv('submission.csv', index=False)
```

Predict and Submit



```
train_and_pred(train_features, test_features, train_labels, test_data,  
               num_epochs, lr, weight_decay, batch_size)
```



Predict and Submit

Step 1

Upload submission file



Upload Submission File

File Format

Your submission should be in CSV format.
You can upload this in a zip/gz/rar/7z archive, if you prefer.

Number of Predictions

We expect the solution file to have 1459 prediction rows. This file should have a header row. Please see sample submission file on the [data page](#).

Step 2

Describe submission

B *I* | 🔗 “ </> 🖼️ | ☰ ☷ H ⚡ | ↺ ↻

 Styling with Markdown supported

Briefly describe your submission.

Make Submission

Thank You !

Does anyone have any questions?

Twitter: @walkercet

Blog: <https://ceteongvanness.wordpress.com>

Resources

Dive into Deep Learning